

DEVELOPMENT: SPA-Net 2M vs. 10M -- classification loss  
(training-time cross-entropy, not AUC)

